

Introduction to Pandas

Lesson 3: Data Manipulation & Processing

An  Commons initiative



<https://creativecommons.org/licenses/by-nc-nd/4.0/>



What is Pandas?

- Officially stands for 'Python Data Analysis Library'
- Powerful Python library for data manipulation and processing
- Your data's home



Pandas install + import

```
conda install pandas or pip install pandas ← from terminal / command line
```

```
!pip install pandas ← in jupyter notebook / Colab notebook
```

```
import pandas as pd ← or can just import pandas in the Colab
```



Core components of Pandas

Series

	oranges
0	0
1	3
2	7
3	2

Dataframe

	apples	oranges
0	3	0
1	2	3
2	0	7
3	1	2

Creating series from scratch

INPUT: `pd.Series([1, 3, 5, np.nan, 6, 8])`

OUTPUT:

0	1.0
1	3.0
2	5.0
3	NaN
4	6.0
5	8.0

Source: https://pandas.pydata.org/pandas-docs/stable/getting_started/10min.html



Creating dataframes from scratch

INPUT: `data = { 'Amy': [80, 60, 90, 80], 'Ben': [78, 50, 100, 80] }`

```
marks = pd.DataFrame(data)
```

marks

OUT:

	Amy	Ben
0	80	78
1	60	50
2	90	100
3	80	80



Read and export to different formats

Data format	Read	Save
.CSV	<code>pd.read_csv()</code>	<code>pd.to_csv()</code>
json	<code>pd.read_json()</code>	<code>pd.to_json()</code>
Excel	<code>pd.read_excel()</code>	<code>pd.to_excel()</code>
SQL	<code>pd.read_sql()</code>	<code>pd.to_sql()</code>



Importing .csv files in Colab

There are different ways of importing .csv files to a colab environment.

For example:

- Your .csv file is on a local drive:

```
[6] 1 import pandas as pd
    2 from google.colab import files
    3 import io
    4
    5 local_drive_file=files.upload()
    6 # you then have to press the 'Choose Files' button, and indicate in the dialog box which file you want to upload (e.g., 'wildlife.csv').
    7 your_data_frame = pd.read_csv(io.BytesIO(local_drive_file['wildlife.csv']))
```

Choose Files wildlife.csv

- **wildlife.csv**(text/csv) - 359 bytes, last modified: 1/20/2021 - 100% done
Saving wildlife.csv to wildlife (1).csv

- Your .csv file is on GitHub:

```
▶ 1 import pandas as pd
  2 # Reading the csv file into a pandas dataframe
  3 # setting the index column as the uniq_id at the same time
  4
  5 wildlife_df = pd.read_csv('https://raw.githubusercontent.com/AICP-teaching/Datasets/main/wildlife.csv')
  6
```



Main questions data scientists face

- How do we view or get info from data?
- What if there are duplicate rows?
- Why and how to change the column names?
- What do I do with missing values?
- What do I do if I just want part of the data frame?



How to view data

head() shows the beginning of a file

```
movies_df.head()
```

tail() shows the end of a file

```
movies_df.tail()
```

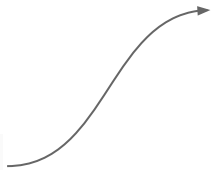
y	0	a	0.4
u	0	b	1.5
y	1	c	6.3
x	0	d	5.3
i	0	e	9.2
...
e	1	h	1.3
t	0	h	4.5
y	0	w	5.3
t	1	d	1.4
i	0	e	2.3



Another viewing option in Pandas

```
>>> pandas.set_option('display.max_rows', 10)
```

```
print(movies_df)
```



y	0	a	0.4
u	0	b	1.5
y	1	c	6.3
x	0	d	5.3
i	0	e	9.2
...
e	1	h	1.3
t	0	h	4.5
y	0	w	5.3
t	1	d	1.4
i	0	e	2.3



Getting info on the data

```
movie_df.info()
```

OUTPUT:

```
<class 'pandas.core.frame.DataFrame'>
Index: 1000 entries, Guardians of the Galaxy to Nine Lives
Data columns (total 11 columns):
Rank                1000 non-null int64
Genre               1000 non-null object
Description         1000 non-null object
Director            1000 non-null object
Actors              1000 non-null object
Year                1000 non-null int64
Runtime (Minutes)   1000 non-null int64
Rating              1000 non-null float64
Votes              1000 non-null int64
Revenue (Millions)  872 non-null float64
Metascore           936 non-null float64
dtypes: float64(3), int64(4), object(4)
memory usage: 93.8+ KB
```



How to know the shape of the data

```
1 wildlife_df
```

	animal	uniq_id	water_need
0	elephant	1001	500
1	elephant	1002	600
2	elephant	1003	550
3	tiger	1004	300
4	tiger	1005	320
5	tiger	1006	330
6	tiger	1007	290
7	tiger	1008	310
8	zebra	1009	200
9	zebra	1010	220
10	zebra	1011	240
11	zebra	1012	230
12	zebra	1013	220
13	zebra	1014	100
14	zebra	1015	80
15	lion	1016	420
16	lion	1017	600
17	lion	1018	500
18	lion	1019	390
19	hyppo	1020	410
20	hyppo	1021	430
21	hyppo	1022	410

(#row, #col)

22

3

```
1 wildlife_df.shape
2
(22, 3)
```

Data's statistical overview

```
1 wildlife_df
```

	animal	uniq_id	water_need
0	elephant	1001	500
1	elephant	1002	600
2	elephant	1003	550
3	tiger	1004	300
4	tiger	1005	320
5	tiger	1006	330
6	tiger	1007	290
7	tiger	1008	310
8	zebra	1009	200
9	zebra	1010	220
10	zebra	1011	240
11	zebra	1012	230
12	zebra	1013	220
13	zebra	1014	100
14	zebra	1015	80
15	lion	1016	420
16	lion	1017	600
17	lion	1018	500
18	lion	1019	390
19	hyppo	1020	410
20	hyppo	1021	430
21	hyppo	1022	410

for all data fields

```
1 wildlife_df.describe()
```

	uniq_id	water_need
count	22.000000	22.000000
mean	1011.500000	347.727273
std	6.493587	147.549243
min	1001.000000	80.000000
25%	1006.250000	232.500000
50%	1011.500000	325.000000
75%	1016.750000	427.500000
max	1022.000000	600.000000

for specific data fields

```
1 wildlife_df['water_need'].describe()
```

count	22.000000
mean	347.727273
std	147.549243
min	80.000000
25%	232.500000
50%	325.000000
75%	427.500000
max	600.000000
Name: water_need, dtype: float64	

describe() more in depth

df.describe()

	Age	Score
count	12.000000	12.000000
mean	32.500000	73.000000
std	9.209679	17.653225
min	24.000000	44.000000
25%	25.750000	64.000000
50%	29.000000	74.000000
75%	35.250000	87.500000
max	51.000000	99.000000

df.describe(include = "all")

	Age	Name	Score
count	12.000000	12	12.000000
unique	NaN	12	NaN
top	NaN	Rahul	NaN
freq	NaN	1	NaN
mean	32.500000	NaN	73.000000
std	9.209679	NaN	17.653225
min	24.000000	NaN	44.000000
25%	25.750000	NaN	64.000000
50%	29.000000	NaN	74.000000
75%	35.250000	NaN	87.500000
max	51.000000	NaN	99.000000



<https://images.app.goo.gl/byfMyVrWQKqEsFNk9>

<https://images.app.goo.gl/6iiMJu9f9ohzVHck15>

Handling duplicates

Duplicates are managed using `drop_duplicates()`

```
movie_df.drop_duplicates( inplace = True )
```

movie_df

*	*
*	*
*	*



movie_df

*	*
*	*
*	*

drop_duplicates() options

```
movie_df.drop_duplicates(keep = ???)
```

- **first**: Drop duplicates except for the first occurrence. **(default)**
- **last**: Drop duplicates except for the last occurrence.
- **False**: Drop all duplicates.

Movie_df - **first**

*	*
*	*
*	*

Movie_df - **last**

*	*
*	*
*	*

Movie_df - **False**

*	*
*	*
*	*



Get columns names with .columns

IN:

	Popularity	Movie_type	Screening
Title			
Kati kati	2	Drama	2016
The letter	1	Documentary	2019

.columns

OUT: Index(['Popularity', 'Movie_type', 'Screening'], dtype='object')



Rename columns

```
1. movies_df.rename(columns={
    'Popularity': 'Rank',
    'Movie_type': 'Genre'
    'Screening': 'Year' },inplace=True)
```

Or ...

```
2. movies_df.columns = ['Rank', 'Genre', 'Year']
```

	Rank	Genre	Year
Title			
Kati kati	2	Drama	2016
The letter	1	Documentary	2019



Using list comprehension

You can iterate over an **iterable**, to modify a **member** object using an **expression**


`movies_df.columns = [col.lower() for col in movies_df]`

	rank	genre	year
Title			
Kati kati	2	Drama	2016
The letter	1	Documentary	2019



Identify missing values

Where are the missing values?

```
movies_df.isnull().sum()
```

	rank	genre	year
Title			
Kati kati	2		2016
The letter	1	Documentary	

	rank	genre	year
Title			
Kati kati	False	True	False
The letter	False	False	True

```
rank      0
genre     1
year      1
dtype: int64
```

.isnull()



.sum()



Removing missing values in rows

Before .dropna():

.dropna()



After .dropna():

	rank	genre	year
Title			
Kati kati	2		
The letter	1	Documentary	2019

	rank	genre	year
Title			
The letter	1	Documentary	2019

.dropna() is equal to .dropna(axis=0). 'axis' refers to the shape: 0=row, 1=column



Removing missing values in columns **Learn**

Before `.dropna(axis = 1):`

`.dropna(axis=1)`

	rank	genre	year
Title			
Kati kati	2		
The letter	1	Documentary	2019



After `.dropna(axis = 1):`

	rank
Title	
Kati kati	2
The letter	1



Filling missing values (Imputation)

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	
Soul boy	4	Drama	2010	1.2

`col.mean()` = 5.23

`col`

```
col = movie_df['revenues_m$']  
mean=col.mean()  
col.fillna(mean, inplace = True)
```

```
col.fillna(movie_df['revenues_m$'].mean(), inplace=True) ➡
```

revenues_m\$
5.6
8.9
5.23
1.2



Series and dataframes

```
movies_df['genre']
```

```
type: series
```

```
movies_df[['genre']]
```

```
type: dataframe
```

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	4.3
Soul boy	4	Drama	2010	1.2



Extracting columns

```
movies_df[['genre', 'year']]  
type: dataframe
```

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	4.3
Soul boy	4	Drama	2010	1.2



	genre	year
Title		
Kati kati	Drama	2016
The letter	Documentary	2019
Nairobi half-life	Drama	2012
Soul boy	Drama	2010



Extract a row by name

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	4.3
Soul boy	4	Drama	2010	1.2



```
movie_df.loc['The letter']
```

	rank	genre	year	revenues_m\$
Title				
The letter	1	Documentary	2019	8.9



Extract a row by index

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	4.3
Soul boy	4	Drama	2010	1.2



```
movie_df.iloc[1]
```

	rank	genre	year	revenues_m\$
Title				
The letter	1	Documentary	2019	8.9



Extract multiple rows

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9
Nairobi half-life	3	Drama	2012	4.3

`movie_df.iloc[:1]`

`movie_df.loc['Kati kati', 'The letter']`

	rank	genre	year	revenues_m\$
Title				
Kati kati	2	Drama	2016	5.6
The letter	1	Documentary	2019	8.9

Conditional selections dataframes

```
[4] 1 wildlife_df
```

	animal	uniq_id	water_need
0	elephant	1001	500
1	elephant	1002	600
2	elephant	1003	550
3	tiger	1004	300
4	tiger	1005	320
5	tiger	1006	330
6	tiger	1007	290
7	tiger	1008	310
8	zebra	1009	200
9	zebra	1010	220
10	zebra	1011	240
11	zebra	1012	230
12	zebra	1013	220
13	zebra	1014	100
14	zebra	1015	80
15	lion	1016	420
16	lion	1017	600
17	lion	1018	500
18	lion	1019	390
19	hyppo	1020	410
20	hyppo	1021	430
21	hyppo	1022	410

```
1 condition=(wildlife_df["animal"]=="zebra")  
2 condition
```


```
0      False  
1      False  
2      False  
3      False  
4      False  
5      False  
6      False  
7      False  
8       True  
9       True  
10      True  
11      True  
12      True  
13      True  
14      True  
15     False  
16     False  
17     False  
18     False  
19     False  
20     False  
21     False  
Name: animal, dtype: bool
```

```
1 zebras_df=wildlife_df[wildlife_df['animal']=="zebra"]  
2 zebras_df
```

	animal	uniq_id	water_need
8	zebra	1009	200
9	zebra	1010	220
10	zebra	1011	240
11	zebra	1012	230
12	zebra	1013	220
13	zebra	1014	100
14	zebra	1015	80

Data Transformation

To make the data normally distributed and make them able to meet the assumptions of parametric statistical tests.



The Effect of Log Transformation. Source: https://www.medcalc.org/manual/log_transformation.php



Time Data

Convert time data to a usable format for manipulations.

```
names = pd.Series(['Daniel', 'Joseph', 'James'])
birthdays = pd.Series(["4th of July, 2015", "19th of Oct, 2015",
                       "3rd of Sept, 2012"])
data = pd.DataFrame({ 'Name': names, 'Birthday': birthdays })
data["DOB"] = pd.to_datetime(data["Birthday"])
data
```

	Name	Birthday	DOB
0	Daniel	4th of July, 2015	2015-07-04
1	Joseph	19th of Oct, 2015	2015-10-19
2	James	3rd of Sept, 2012	2012-09-03

```
data['year'] = data["DOB"].dt.year
data[data['year']==2015]['Name']
```

```
0    Daniel
1    Joseph
Name: Name, dtype: object
```



Questions?

Thank you for your attention!



References and further readings

- <https://pythongeeks.net/python-tutorials/python-pandas-for-beginners-a-complete-guide/>
- <https://towardsdatascience.com/python-for-data-science-basics-of-pandas-5f8d9680617e>
- https://www.learndatasci.com/tutorials/python-pandas-tutorial-complete-introduction-for-beginners/https://pandas.pydata.org/pandas-docs/stable/getting_started/10min.html
- <https://www.shanelynn.ie/select-pandas-dataframe-rows-and-columns-using-iloc-loc-and-ix/>
- <https://www.dataquest.io/blog/pandas-cheat-sheet/>



Contributors



Elias Williams



Jose Ignacio Diaz



Pavithra Rajasekar



Gilles Fayad

